

Investigate_a_Dataset

November 11, 2020

1 Project: Investigate a Dataset (TMDB movies)

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction - Which year has the highest release of movies? - What are the top 20 most expensive movies? - which Movies with Highest And Lowest Votes? - How do ratings correlate with commercial success (profits)? - What is the average runtime of all movies? ### TMDB movies

This data set contains information about 10,000 movies collected from The Movie Database (TMDB), including user ratings and revenue.

```
In [1]: # Use this cell to set up import statements for all of the packages that you
#       plan to use.
#       Remember to include a 'magic word' so that your visualizations are plotted
#       inline with the notebook. See this page for more:
#       http://ipython.readthedocs.io/en/stable/interactive/magics.html
import numpy as np
import pandas as pd
import matplotlib as plt
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

Data Wrangling

1.1.1 Gathering Data

```
In [2]: # Load your data and print out a few lines. Perform operations to inspect data
#       types and look for instances of missing or possibly errant data.
```

```
data=pd.read_csv ("tmdb-movies.csv")
data.head()
```

```
Out[2]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	

1	76341	tt1392190	28.419936	150000000	378436354
2	262500	tt2908446	13.112507	110000000	295238201
3	140607	tt2488496	11.173104	200000000	2068178225
4	168259	tt2820852	9.335014	190000000	1506249360

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	

	cast	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	Shailene Woodley Theo James Kate Winslet Ansel...	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	

	homepage	director	\
0	http://www.jurassicworld.com/	Colin Trevorrow	
1	http://www.madmaxmovie.com/	George Miller	
2	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke	
3	http://www.starwars.com/films/star-wars-episod...	J.J. Abrams	
4	http://www.furious7.com/	James Wan	

	tagline	...	\
0	The park is open.	...	
1	What a Lovely Day.	...	
2	One Choice Can Destroy You	...	
3	Every generation has a story.	...	
4	Vengeance Hits Home	...	

	overview	runtime	\
0	Twenty-two years after the events of Jurassic ...	124	
1	An apocalyptic story set in the furthest reach...	120	
2	Beatrice Prior must confront her inner demons ...	119	
3	Thirty years after defeating the Galactic Empi...	136	
4	Deckard Shaw seeks revenge against Dominic Tor...	137	

	genres	\
0	Action Adventure Science Fiction Thriller	
1	Action Adventure Science Fiction Thriller	
2	Adventure Science Fiction Thriller	
3	Action Adventure Science Fiction Fantasy	
4	Action Crime Thriller	

	production_companies	release_date	vote_count	\
--	----------------------	--------------	------------	---

0	Universal Studios Amblin Entertainment Legenda...	6/9/15	5562
1	Village Roadshow Pictures Kennedy Miller Produ...	5/13/15	6185
2	Summit Entertainment Mandeville Films Red Wago...	3/18/15	2480
3	Lucasfilm Truenorth Productions Bad Robot	12/15/15	5292
4	Universal Pictures Original Film Media Rights ...	4/1/15	2947

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
2	6.3	2015	1.012000e+08	2.716190e+08
3	7.5	2015	1.839999e+08	1.902723e+09
4	7.3	2015	1.747999e+08	1.385749e+09

[5 rows x 21 columns]

1.1.2 Assess Data

In [3]: data.shape

Out[3]: (10866, 21)

In [4]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64
vote_average      10866 non-null float64
release_year      10866 non-null int64
budget_adj        10866 non-null float64
revenue_adj       10866 non-null float64
dtypes: float64(4), int64(6), object(11)
```

memory usage: 1.7+ MB

In [5]: data.describe()

```
Out[5]:
```

	id	popularity	budget	revenue	runtime \
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

In [6]: data.budget.value_counts()

```
Out[6]:
```

0	5696
20000000	190
15000000	183
25000000	178
10000000	176
30000000	165
5000000	141
40000000	134
35000000	128
12000000	120
50000000	112
6000000	109
8000000	102
3000000	101
60000000	99
7000000	92
4000000	84
2000000	81
18000000	74
1000000	73
70000000	66
11000000	65

45000000	64
80000000	62
13000000	59
17000000	59
14000000	56
22000000	54
3500000	52
75000000	51
...	
9100000	1
34200000	1
818418	1
14200000	1
8400000	1
22997992	1
7000	1
220000	1
23600000	1
1645000	1
82500000	1
27220000	1
5112027	1
75	1
1052753	1
225000	1
115	1
12000	1
163000000	1
2240000	1
786675	1
4250000	1
207000000	1
61733	1
19885552	1
51500000	1
25500000	1
1350000	1
7920000	1
4653000	1

Name: budget, Length: 557, dtype: int64

In [7]: data.revenue.value_counts()

Out[7]:

0	6016
12000000	10
10000000	8
11000000	7
6000000	6

5000000	6
2000000	6
13000000	5
20000000	5
14000000	5
30000000	5
7000000	4
4300000	4
4000000	4
1000000	4
50000000	4
25000000	4
60000000	3
6700000	3
11	3
16000000	3
15	3
29000000	3
21000000	3
100000000	3
102000000	3
3	3
16	3
2500000	3
70000000	3
...	
42721196	1
79958599	1
40084041	1
442965	1
50752337	1
30857814	1
1400000	1
7027290	1
303788635	1
10300000	1
45916769	1
26049082	1
211989043	1
41009669	1
346079773	1
4729352	1
23159305	1
272742922	1
525	1
2626800	1
39946780	1
3117985	1

```
22270      1
65884703    1
4500000     1
53676580    1
617000     1
13001257    1
504050219   1
20518905    1
Name: revenue, Length: 4702, dtype: int64
```

```
In [8]: data.runtime.value_counts()
```

```
Out[8]: 90      547
95      358
100     335
93      328
97      306
96      300
91      297
94      292
92      270
98      270
88      270
89      253
105     250
101     232
99      231
102     228
85      227
103     212
87      211
106     211
108     205
104     201
86      197
107     194
110     194
112     152
84      142
120     140
111     140
109     140
...
294      1
366      1
470      1
550      1
566      1
```

31	1
500	1
372	1
292	1
252	1
257	1
417	1
561	1
705	1
34	1
194	1
210	1
226	1
242	1
250	1
282	1
338	1
19	1
51	1
223	1
235	1
20	1
220	1
236	1
219	1

Name: runtime, Length: 247, dtype: int64

1.1.3 Tidiness Issues

- Drop duplicate Rows.

1.1.4 Quality

- Drop unnecessary columns:imdb_id,homepage,keywords,tag_line,overview,revenue_adj,budget_adj.
- Wrong data types (id need to convert to string).
- Wrong data types (relase_date need to convert to datetime).
- Drop the movies which are having zero value of budget and revenue.
- Drop the movies which having missing data (cast,director,geners,production_companies).

1.1.5 Data Cleaning

Tidiness Issues

Define:

- Removing duplicates rows from dataset.

Code:

```
In [9]: data.drop_duplicates(inplace=True)
```

Test:

```
In [10]: data.shape
```

```
Out[10]: (10865, 21)
```

Quality**Define:**

- removing unnecessary columns:imdb_id,homepage,keywords,tag_line,overview,revenue_adj,budget_adj.

Code:

```
In [11]: data.drop(['imdb_id','homepage','keywords','tagline','overview','revenue_adj','budget_a
```

Test:

```
In [12]: data.shape
```

```
Out[12]: (10865, 14)
```

Define:

- converting id to string.

Code:

```
In [13]: data['id']=str(data['id'])
```

Test:

```
In [14]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 14 columns):
id                10865 non-null object
popularity        10865 non-null float64
budget            10865 non-null int64
revenue           10865 non-null int64
original_title    10865 non-null object
cast              10789 non-null object
director          10821 non-null object
runtime           10865 non-null int64
genres            10842 non-null object
```

```
production_companies    9835 non-null object
release_date            10865 non-null object
vote_count              10865 non-null int64
vote_average            10865 non-null float64
release_year            10865 non-null int64
dtypes: float64(2), int64(5), object(7)
memory usage: 1.2+ MB
```

Define:

- converting release_date to datetime

Code:

```
In [15]: data['release_date']=pd.to_datetime(data['release_date'])
```

Test:

```
In [16]: data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 14 columns):
id                10865 non-null object
popularity        10865 non-null float64
budget           10865 non-null int64
revenue          10865 non-null int64
original_title    10865 non-null object
cast              10789 non-null object
director          10821 non-null object
runtime          10865 non-null int64
genres            10842 non-null object
production_companies 9835 non-null object
release_date      10865 non-null datetime64[ns]
vote_count        10865 non-null int64
vote_average      10865 non-null float64
release_year      10865 non-null int64
dtypes: datetime64[ns](1), float64(2), int64(5), object(6)
memory usage: 1.2+ MB
```

Define:

- removing movies with zero values in budget and revenue.

Code:

```
In [17]: data=data[data['budget']!=0]
         data=data[data['revenue']!=0]
```

Test:

```
In [18]: data.shape
```

```
Out[18]: (3854, 14)
```

Define:

- removing missing data in cast,director,genres,production_companies.

Code:

```
In [19]: data=data[data.cast.notnull()]
         data=data[data.director.notnull()]
         data=data[data.production_companies.notnull()]
```

Test:

```
In [20]: data.shape
```

```
Out[20]: (3805, 14)
```

```
In [21]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3805 entries, 0 to 10848
Data columns (total 14 columns):
id                3805 non-null object
popularity        3805 non-null float64
budget           3805 non-null int64
revenue          3805 non-null int64
original_title    3805 non-null object
cast             3805 non-null object
director         3805 non-null object
runtime          3805 non-null int64
genres           3805 non-null object
production_companies 3805 non-null object
release_date      3805 non-null datetime64[ns]
vote_count        3805 non-null int64
vote_average      3805 non-null float64
release_year      3805 non-null int64
dtypes: datetime64[ns](1), float64(2), int64(5), object(6)
memory usage: 445.9+ KB
```


Exploratory Data Analysis

Tip: Now that you've trimmed and cleaned your data, you're ready to move on to exploration. Compute statistics and create visualizations with the goal of addressing the research questions that you posed in the Introduction section. It is recommended that you be systematic with your approach. Look at one variable at a time, and then follow it up by looking at relationships between variables.

1.1.6 Research Question 1 (Which year has the highest release of movies?)

```
In [22]: data.groupby('release_year').count()['id'].plot(xticks = np.arange(1960,2016,5))

#set the figure size and labels
sns.set(rc={'figure.figsize':(10,5)})
plt.title("Year Vs Number Of Movies",fontsize = 14)
plt.xlabel('Release year',fontsize = 13)
plt.ylabel('Number Of Movies',fontsize = 13)
#set the style sheet
sns.set_style("whitegrid")
```




Movies number increase year by year but after 2010 it start decreases, by calculations that year 2010 has the highest release of movies.

1.1.7 Research Question 2 (What are the top 20 most expensive movies?)

```
In [23]: sorted_budget = data['budget'].sort_values(ascending=False)[:20]
high_budget=pd.DataFrame()
titles_exp=[]
budgets=[]
for i in sorted_budget.index:
    titles_exp.append(data.loc[i,'original_title'])
```

```
        budgets.append(sorted_budget.loc[i])
high_budget['Titles']=titles_exp
high_budget['Budgets']=budgets
high_budget.set_index('Titles',inplace=True)
high_budget.plot(kind='bar',figsize=(8,8))
plt.title('Top 20 most expensive movies (1960 - 2015) ');
plt.ylabel('Budget in 100\'s of million ($)');
```




The top 20 most expensive movies ("The Warrior's Way", 'Pirates of the Caribbean: On Stranger Tides', "Pirates of the Caribbean: At World's End", 'Avengers: Age of Ultron', 'Superman Returns', 'Tangled', 'John Carter', 'Spider-Man 3', 'The Lone Ranger', 'Harry Potter and the Deathly Hallows: Part 1', 'Harry Potter and the Half-Blood Prince', 'The Dark Knight Rises', 'X-Men: Days of Future Past', 'The Hobbit: The Battle of the Five Armies', 'The Hobbit: The Desolation of Smaug', 'The Hobbit: An Unexpected Journey', 'Spectre', 'Avatar', 'Man of Steel', 'Men in Black 3').

1.1.8 Research Question 3 (which Movies with Highest And Lowest Votes?)

```
In [24]: info = pd.DataFrame(data['vote_average'].sort_values(ascending = False))
        info['original_title'] = data['original_title']
        d = list(map(str, (info['original_title'])))

        ##extract the top 10 highly rated movies data from the list and dataframe.
        x = list(d[:10])
        y = list(info['vote_average'][:10])


        #make the point plot and setup the title and labels.
        ax = sns.pointplot(x=y,y=x)
        sns.set(rc={'figure.figsize':(10,5)})
        ax.set_title("Top 10 Highest Rated Movies",fontsize = 15)
        ax.set_xlabel("Vote Average",fontsize = 13)
        #setup the stylesheet
        sns.set_style("darkgrid")
```



The top highest rated movies ('Stop Making Sense', 'The Shawshank Redemption', 'The Godfather', 'Whiplash', 'Pulp Fiction', 'The Dark Knight', 'Forrest Gump', 'The Godfather: Part II', 'Schindler's List', 'Fight Club').

1.1.9 Research Question 4 (How do ratings correlate with commercial success (profits)?)

```
In [25]: d=data
d['Profit']=d['revenue']-d['budget']
d.plot(x='vote_average',y='Profit',kind='scatter',figsize=(8,8));
plt.ylabel('Profit in billion ($)');
plt.xlabel('Rating');
plt.title('Rating vs Profit');
```



Movies with high rating have the highest profit.

Research Question 5 (What is the average runtime of all movies?)


```
In [26]: def average_func(column_name):  
  
         return data[column_name].mean()  
         average_func('runtime')
```

```
Out[26]: 109.35111695137977
```

The average runtime of all movies in this dataset is 109 mins approx. ## Conclusions

1.2 Submitting your Project

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File > Download as** sub-menu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

1.2.1 Conclusions

- 2010 year has the highest release of movies.
- "The Warrior's Way" is the most expensive movie.
- "Stop Making Sense" is the highest rating movie.
- The average runtime of all movies in this dataset is 109 mins approx.

1.2.2 Limitation

- Units of revenue and budget column: I am not sure that the budgets and revenues all in US dollars?
- the results can only be treated as indicators and are not generalizable.

```
In [27]: from subprocess import call  
         call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

```
Out[27]: 0
```

```
In [ ]:
```

```
In [ ]:
```